

Mini Project on CO2 Emission Data Using PySpark

Dataset : CO2 Emissions_Canada.csv

Source : <https://www.kaggle.com/datasets/debajyotipodder/co2-emission-by-vehicles>

```
import os
import findspark

# Point to Java & Spark
os.environ["JAVA_HOME"] = "C:/Progra~1/Java/jdk1.8"
os.environ["SPARK_HOME"] = "C:/spark/spark-3.5.7-bin-hadoop3-
scala2.13"
os.environ["HADOOP_HOME"] = "C:/hadoop" # if you installed winutils
here
os.environ["PATH"] +=
";C:/spark/spark-3.5.7-bin-hadoop3-scala2.13/bin;C:/hadoop/bin"

# Initialize findspark
findspark.init(os.environ["SPARK_HOME"])

from pyspark.sql import SparkSession

# Now build SparkSession
spark = SparkSession.builder \
    .appName("BDA_MINI_PROJECT") \
    .master("local[*]") \
    .getOrCreate()

sc = spark.sparkContext

sc

<SparkContext master=local[*] appName=BDA_MINI_PROJECT>
```

Dataset Insights

The dataset provided contains details of 7,385 vehicles across different makes, models, classes, and fuel types in Canada. It has 12 columns: Make, Model, Vehicle Class, Engine Size(L), Cylinders, Transmission, Fuel Type, Fuel Consumption (City, Hwy, Comb), Fuel Consumption Comb (mpg), and CO2 Emissions(g/km). There are no missing values, ensuring data consistency and completeness, which is advantageous for analysis. Below are the detailed insights:

1. Structure and Uniqueness

The dataset comprises 7,385 rows and 12 columns. Each vehicle is uniquely identified by the combination of Make and Model. There are 42 unique makes (e.g., Ford, Toyota, BMW), and over 2,000 unique models. The dataset captures a wide range of vehicle types, including 16 vehicle classes (from two-seaters to large SUVs and pickups). This diversity provides a strong foundation for analyzing the automotive market in terms of fuel efficiency and emissions.

1. Vehicle Class Distribution

The dataset spans 16 vehicle classes, such as Compact, Midsize, SUV, Pickup Truck, and Station Wagon. SUVs and Pickup Trucks appear with higher representation, reflecting their popularity in the Canadian market. Smaller classes like Two-Seater and Subcompact also exist but in smaller numbers. Such distribution helps analyze the impact of vehicle type on CO2 emissions and fuel economy.

1. Transmission and Engine Configurations

There are 27 unique transmission types ranging from manual to automatic with varying gears, and 8 unique cylinder configurations (from 3 to 12 cylinders). Vehicles with higher cylinder counts generally correspond to larger engines and higher emissions. This allows comparative analysis of drivetrain efficiency.

1. CO2 Emissions Insights

The CO2 Emissions(g/km) variable ranges widely, with values as low as ~100 g/km for efficient hybrids to well above 400 g/km for large, less efficient vehicles. The average CO2 emission across all vehicles is around ~250 g/km, showing significant variation. This indicates the presence of both eco-friendly cars and high-emission performance vehicles.

1. Fuel Consumption Distribution

Fuel consumption is provided for city, highway, and combined driving in liters per 100 km. City consumption values are generally higher than highway, reflecting real-world driving conditions. The combined fuel consumption ranges from ~4 L/100 km (very efficient hybrids) to over 20 L/100 km (large SUVs and pickups), indicating major differences in fuel economy across vehicle classes.

1. Fuel Type Patterns

The dataset includes 5 fuel types: Regular gasoline, Premium gasoline, Diesel, Ethanol (E85), and Natural Gas. Regular and Premium gasoline vehicles dominate the dataset, while Diesel and alternative fuels have a smaller share. Average CO2 emissions vary by fuel type, with gasoline vehicles contributing higher emissions on average compared to hybrid/alternative fuel options.

1. Key Relationships

Engine Size vs. CO2: Larger engine sizes (L) generally correlate with higher CO2 emissions.

Cylinders vs. Fuel Consumption: Vehicles with more cylinders show higher city and highway fuel consumption.

Transmission vs. Efficiency: Vehicles with newer automatic transmissions often achieve better combined fuel economy compared to older manual configurations.

Fuel Type vs. Emissions: Alternative fuels (E85, Hybrid, CNG) generally report lower CO2 than traditional gasoline.

1. Market Characteristics

Overall, the dataset portrays a diverse Canadian automotive market with a wide mix of vehicles across size, power, and efficiency. The absence of missing data makes this dataset reliable for

environmental policy analysis, automotive research, and predictive modeling of emissions. It highlights key industry trends: strong SUV and pickup representation, dominance of gasoline fuel types, and the growing presence of fuel-efficient models.

```
import pandas as pd
import matplotlib.pyplot as plt
```

```
# Load dataset
```

```
df = pd.read_csv("C02 Emissions_Canada.csv")
df
```

	Make	Model	Vehicle Class	Engine Size(L)	Cylinders	\
0	ACURA	ILX	COMPACT	2.0	4	
1	ACURA	ILX	COMPACT	2.4	4	
2	ACURA	ILX HYBRID	COMPACT	1.5	4	
3	ACURA	MDX 4WD	SUV - SMALL	3.5	6	
4	ACURA	RDX AWD	SUV - SMALL	3.5	6	
...	
7380	VOLVO	XC40 T5 AWD	SUV - SMALL	2.0	4	
7381	VOLVO	XC60 T5 AWD	SUV - SMALL	2.0	4	
7382	VOLVO	XC60 T6 AWD	SUV - SMALL	2.0	4	
7383	VOLVO	XC90 T5 AWD	SUV - STANDARD	2.0	4	
7384	VOLVO	XC90 T6 AWD	SUV - STANDARD	2.0	4	

	Transmission	Fuel Type	Fuel Consumption City (L/100 km)	\
0	AS5	Z	9.9	
1	M6	Z	11.2	
2	AV7	Z	6.0	
3	AS6	Z	12.7	
4	AS6	Z	12.1	
...	
7380	AS8	Z	10.7	
7381	AS8	Z	11.2	
7382	AS8	Z	11.7	
7383	AS8	Z	11.2	
7384	AS8	Z	12.2	

	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)	\
0	6.7		
8.5			
1	7.7		
9.6			
2	5.8		
5.9			
3	9.1		
11.1			
4	8.7		
10.6			
...

.	
7380	7.7
9.4	
7381	8.3
9.9	
7382	8.6
10.3	
7383	8.3
9.9	
7384	8.7
10.7	

	Fuel Consumption Comb (mpg)	C02 Emissions(g/km)
0	33	196
1	29	221
2	48	136
3	25	255
4	27	244
...
7380	30	219
7381	29	232
7382	27	240
7383	29	232
7384	26	248

[7385 rows x 12 columns]

Dataset Insights

```
print("Dataset Shape:", df.shape)
print("Columns:", df.columns.tolist())
print("Missing Values:\n", df.isnull().sum())
print("Unique values per column:\n", df.nunique())
print(df.describe())
```

Dataset Shape: (7385, 12)

Columns: ['Make', 'Model', 'Vehicle Class', 'Engine Size(L)', 'Cylinders', 'Transmission', 'Fuel Type', 'Fuel Consumption City (L/100 km)', 'Fuel Consumption Hwy (L/100 km)', 'Fuel Consumption Comb (L/100 km)', 'Fuel Consumption Comb (mpg)', 'C02 Emissions(g/km)']

Missing Values:

Make	0
Model	0
Vehicle Class	0
Engine Size(L)	0
Cylinders	0
Transmission	0
Fuel Type	0
Fuel Consumption City (L/100 km)	0
Fuel Consumption Hwy (L/100 km)	0
Fuel Consumption Comb (L/100 km)	0

```

Fuel Consumption Comb (mpg)      0
CO2 Emissions(g/km)              0
dtype: int64
Unique values per column:
  Make                             42
  Model                           2053
  Vehicle Class                    16
  Engine Size(L)                   51
  Cylinders                        8
  Transmission                     27
  Fuel Type                        5
  Fuel Consumption City (L/100 km) 211
  Fuel Consumption Hwy (L/100 km)  143
  Fuel Consumption Comb (L/100 km) 181
  Fuel Consumption Comb (mpg)       54
  CO2 Emissions(g/km)              331
dtype: int64

```

	Engine Size(L)	Cylinders	Fuel Consumption City (L/100 km)
count	7385.000000	7385.000000	7385.000000
mean	3.160068	5.615030	12.556534
std	1.354170	1.828307	3.500274
min	0.900000	3.000000	4.200000
25%	2.000000	4.000000	10.100000
50%	3.000000	6.000000	12.100000
75%	3.700000	6.000000	14.600000
max	8.400000	16.000000	30.600000


	Fuel Consumption Hwy (L/100 km)	Fuel Consumption Comb (L/100 km)
count	7385.000000	7385.000000
mean	9.041706	10.975071
std	2.224456	2.892506
min	4.000000	4.100000
25%	7.500000	8.900000
50%	8.700000	10.600000

75%	10.200000
12.600000	
max	20.600000
26.100000	


	Fuel Consumption Comb (mpg)	CO2 Emissions(g/km)
count	7385.000000	7385.000000
mean	27.481652	250.584699
std	7.231879	58.512679
min	11.000000	96.000000
25%	22.000000	208.000000
50%	27.000000	246.000000
75%	32.000000	288.000000
max	69.000000	522.000000

#1.Distribution of CO2 Emissions

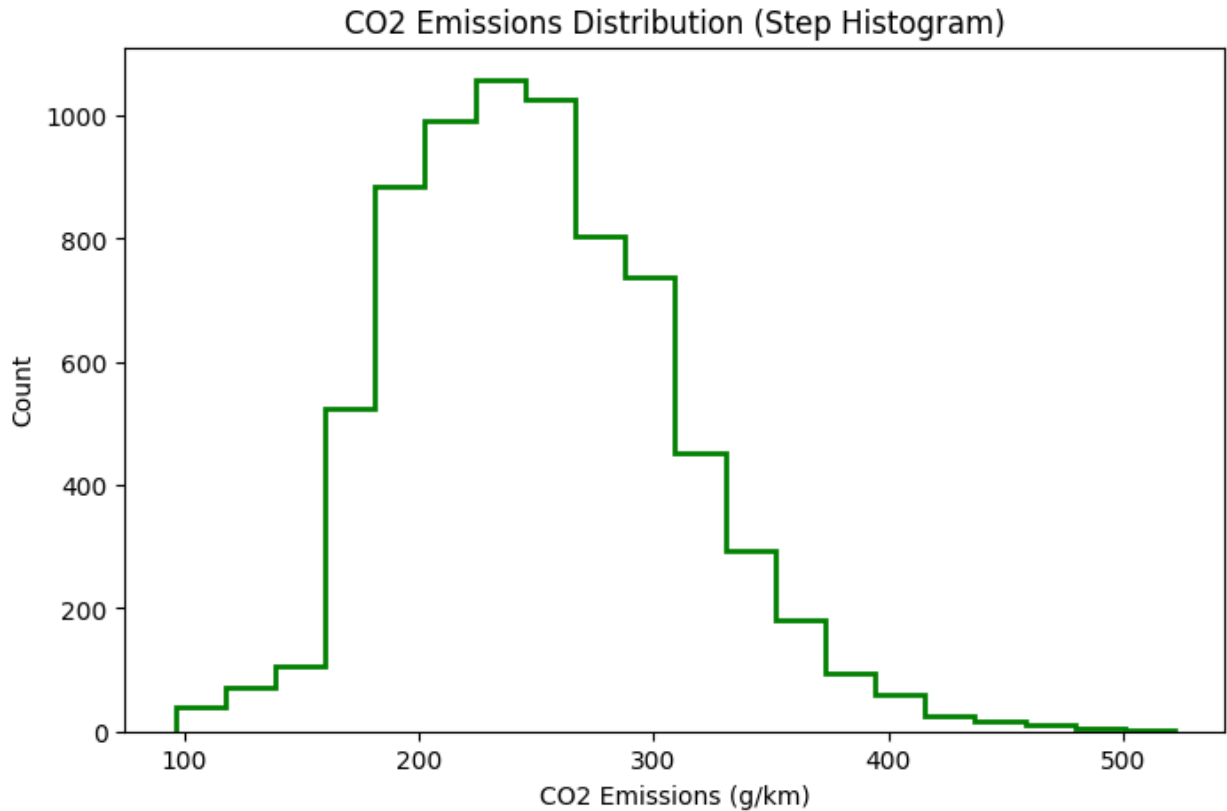
```
plt.figure(figsize=(8,5))
plt.hist(df['CO2 Emissions(g/km)'], bins=20, color="skyblue",
edgecolor="black")
plt.title("Distribution of CO2 Emissions")
plt.xlabel("CO2 Emissions (g/km)")
plt.ylabel("Count")
plt.show()
```



```
#2. Fuel Consumption Distribution (City vs Highway)
plt.figure(figsize=(8,5))
plt.hist(df['Fuel Consumption City (L/100 km)'], bins=20, alpha=0.6,
label="City")
plt.hist(df['Fuel Consumption Hwy (L/100 km)'], bins=20, alpha=0.6,
label="Highway")
plt.title("Fuel Consumption Distribution")
plt.xlabel("Fuel Consumption (L/100 km)")
plt.ylabel("Count")
plt.legend()
plt.show()
```




```
#3. CO2 Emissions Distribution → Step Histogram
plt.figure(figsize=(8,5))
plt.hist(df['CO2 Emissions(g/km)'], bins=20, color="green",
histtype="step", linewidth=2)
plt.title("CO2 Emissions Distribution (Step Histogram)")
plt.xlabel("CO2 Emissions (g/km)")
plt.ylabel("Count")
plt.show()
```




#4. Average CO2 Emissions by Fuel Type

```
avg_co2_fuel = df.groupby("Fuel Type")["CO2  
Emissions(g/km)"].mean().sort_values()  
plt.figure(figsize=(8,5))  
avg_co2_fuel.plot(kind="barh", color="orange", edgecolor="black")  
plt.title("Average CO2 Emissions by Fuel Type")  
plt.xlabel("Average CO2 Emissions (g/km)")  
plt.ylabel("Fuel Type")  
plt.show()
```

#5. Average CO2 Emissions by Fuel Type


```
avg_co2 = df.groupby("Fuel Type")["CO2 Emissions(g/km)"].mean()
plt.figure(figsize=(8,5))
plt.plot(avg_co2.index, avg_co2.values, marker="o", linestyle="--",
color="red")
plt.title("Average CO2 Emissions by Fuel Type")
plt.xlabel("Fuel Type")
plt.ylabel("Average CO2 (g/km)")
plt.xticks(rotation=45)
plt.show()
```




#6. Vehicle Count by Fuel Type

```
fuel_count = df["Fuel Type"].value_counts()
plt.figure(figsize=(6,6))
plt.pie(fuel_count, labels=fuel_count.index, autopct="%1.1f%%",
startangle=90)
plt.title("Vehicle Count by Fuel Type")
plt.show()
```

Vehicle Count by Fuel Type

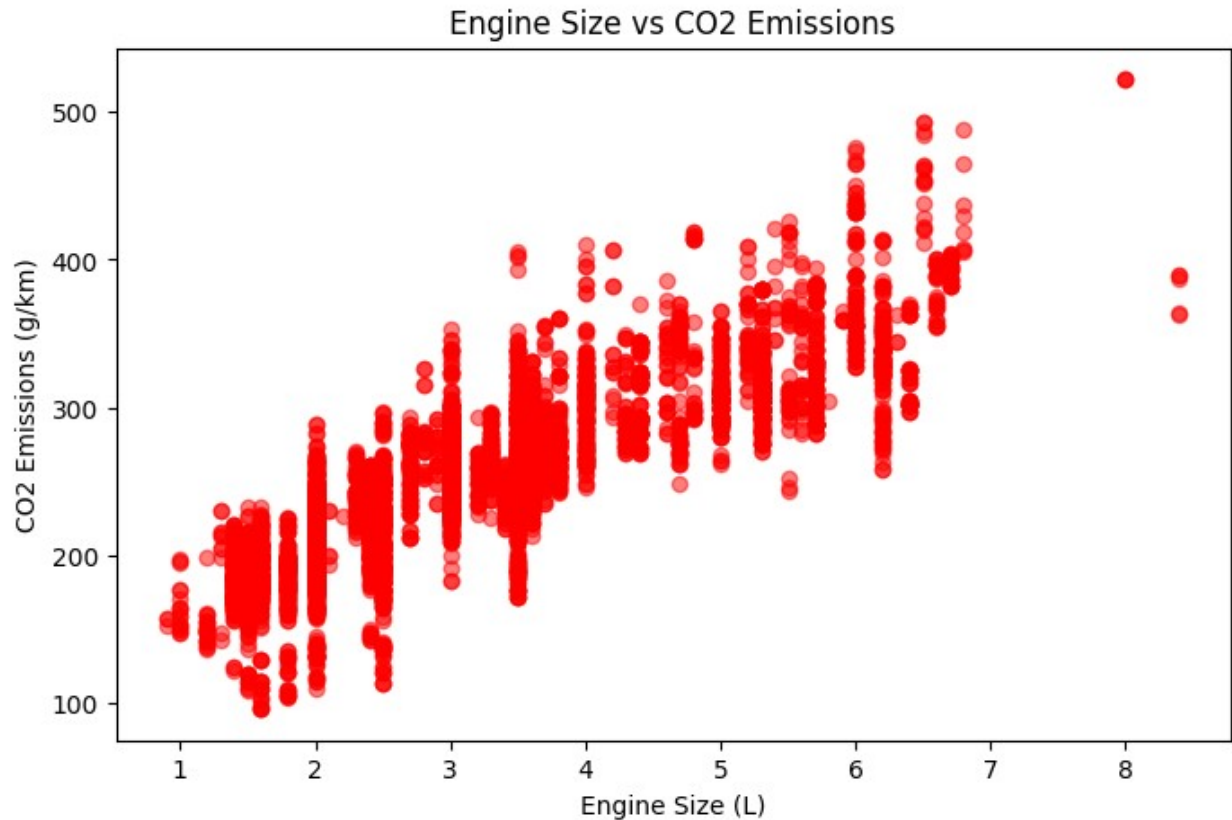


```
#7.Vehicle Count by Vehicle Class
class_count = df["Vehicle Class"].value_counts()
plt.figure(figsize=(8,5))
plt.barh(class_count.index, class_count.values, color="navy")
plt.title("Vehicle Count by Class")
plt.xlabel("Count")
plt.ylabel("Vehicle Class")
plt.show()
```

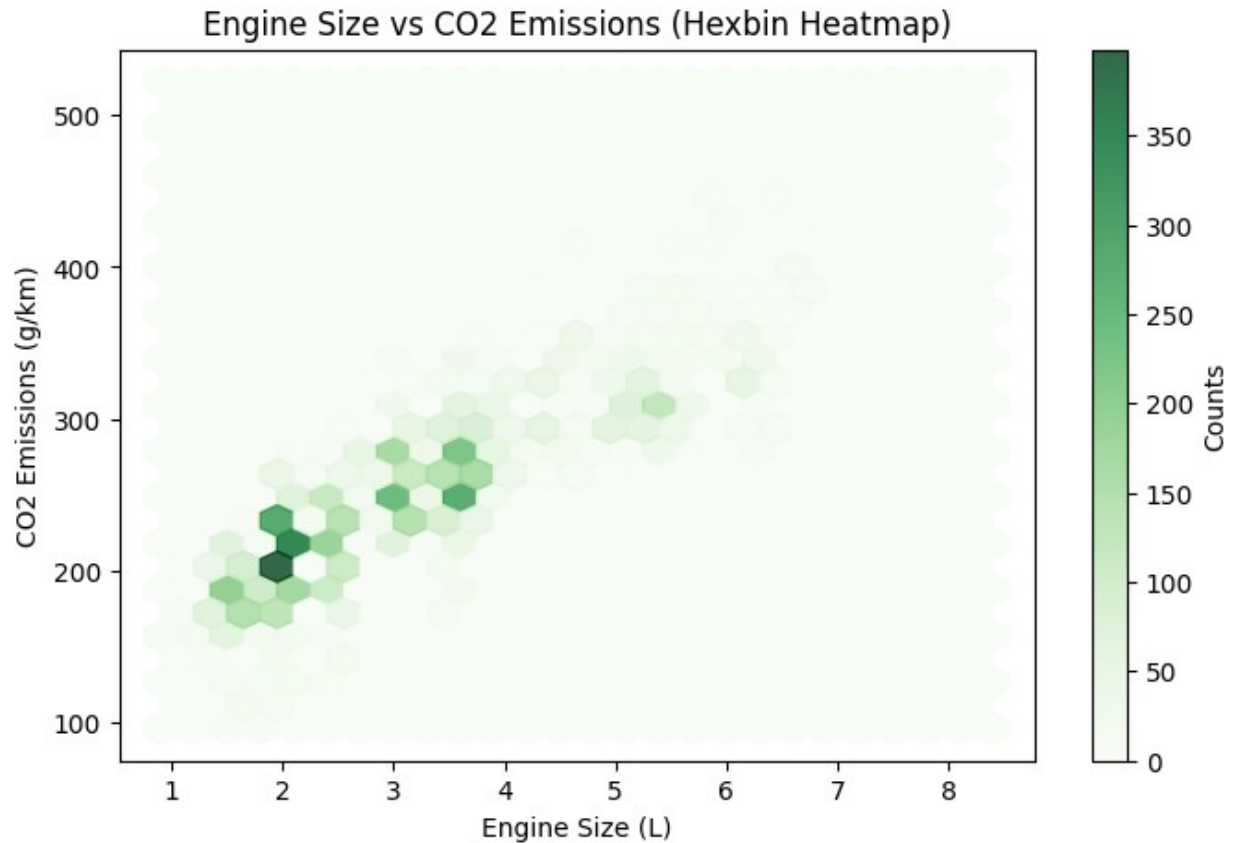


#8. Scatter Plot → Engine Size vs CO2 Emissions

```
plt.figure(figsize=(8,5))
plt.scatter(df['Engine Size(L)'], df['CO2 Emissions(g/km)'],
alpha=0.5, c='red')
plt.title("Engine Size vs CO2 Emissions")
plt.xlabel("Engine Size (L)")
plt.ylabel("CO2 Emissions (g/km)")
plt.show()
```




```
#9.Engine Size vs CO2 Emissions (Hexbin Heatmap)
plt.figure(figsize=(8,5))
plt.hexbin(df['Engine Size(L)'], df['CO2 Emissions(g/km)'],
gridsize=25, cmap="Greens", alpha=0.8)
plt.colorbar(label="Counts")
plt.title("Engine Size vs CO2 Emissions (Hexbin Heatmap)")
plt.xlabel("Engine Size (L)")
plt.ylabel("CO2 Emissions (g/km)")
plt.show()
```



#10. Boxplot → CO2 by Transmission

```
plt.figure(figsize=(10,6))
df.boxplot(column="CO2 Emissions(g/km)", by="Transmission",
grid=False, showfliers=False)
plt.title("CO2 Emissions by Transmission")
plt.suptitle("")
plt.xlabel("Transmission")
plt.ylabel("CO2 Emissions (g/km)")
plt.xticks(rotation=45)
plt.show()
```

<Figure size 1000x600 with 0 Axes>



```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, avg, min, max

# Start Spark session
spark =
SparkSession.builder.appName("CO2EmissionsCanada").getOrCreate()

# Load the dataset
df = spark.read.option("header", True).option("inferSchema",
True).csv("CO2 Emissions_Canada.csv")

#1. Total number of vehicles-----
print("Total number of vehicles:", df.count())

Total number of vehicles: 7385

#2. Total number of fuel types-----
print("Total number of fuel types:", df.select("Fuel
Type").distinct().count())

Total number of fuel types: 5
```

#3. Fuel Type Names-----

```
print("Fuel Types:")
df.select("Fuel Type").distinct().show()
```

Fuel Types:

Fuel Type
E
D
Z
N
X

#4. Total number of vehicles by fuel type-----

```
print("Vehicles by Fuel Type:")
df.groupBy("Fuel Type").count().show()
```

Vehicles by Fuel Type:

Fuel Type	count
E	370
D	175
Z	3202
N	1
X	3637

#5. Total number of vehicles by vehicle class-----

```
print("Vehicles by Vehicle Class:")
df.groupBy("Vehicle Class").count().show()
```

Vehicles by Vehicle Class:

Vehicle Class	count
MINICOMPACT	326
SPECIAL PURPOSE V...	77
COMPACT	1022
VAN - CARGO	22
SUV - STANDARD	735
SUV - SMALL	1217
PICKUP TRUCK - ST...	538
STATION WAGON - M...	53
TWO-SEATER	460
SUBCOMPACT	606
MID-SIZE	1133

PICKUP TRUCK - SMALL	159
FULL-SIZE	639
STATION WAGON - S...	252
MINIVAN	80
VAN - PASSENGER	66

#6. Vehicles grouped by Fuel Type & Transmission----

```
print("Vehicles by Fuel Type and Transmission:")
```

```
df.groupby("Fuel Type", "Transmission").count().show()
```

Vehicles by Fuel Type and Transmission:

Fuel Type	Transmission	count
Z	AV6	9
D	AM6	6
X	A5	76
X	AV10	4
Z	A6	55
E	AS8	18
Z	AS5	4
X	A8	234
Z	M6	465
D	AS10	14
X	AS9	24
Z	AS9	53
Z	AS8	946
Z	AM8	62
Z	A7	53
Z	AV	6
X	AS5	22
D	A6	36
E	A6	166
Z	AV7	30

only showing top 20 rows

#7. Min and Max CO2 emissions per vehicle class----

```
print("Minimum and maximum CO2 emissions in each vehicle class:")
```

```
df.groupby("Vehicle Class").agg(
    min("CO2 Emissions(g/km)").alias("min_co2"),
    max("CO2 Emissions(g/km)").alias("max_co2")
).show()
```

Minimum and maximum CO2 emissions in each vehicle class:

Vehicle Class	min_co2	max_co2
---------------	---------	---------

MINICOMPACT	143	365
SPECIAL PURPOSE V...	208	298
COMPACT	106	404
VAN - CARGO	334	396
SUV - STANDARD	156	476
SUV - SMALL	128	353
PICKUP TRUCK - ST...	238	414
STATION WAGON - M...	126	386
TWO-SEATER	145	522
SUBCOMPACT	152	392
MID-SIZE	104	465
PICKUP TRUCK - SMALL	235	331
FULL-SIZE	96	404
STATION WAGON - S...	110	382
MINIVAN	221	296
VAN - PASSENGER	315	488

#8. Vehicles with C02 > average-----

```
avg_co2_val = df.agg(avg("C02
Emissions(g/km)").alias("avg_co2")).collect()[0]["avg_co2"]
```

```
print("Vehicles with C02 Emissions greater than average:")
df.filter(col("C02 Emissions(g/km)") > avg_co2_val) \
    .select("Make", "Model", "Fuel Type", "C02 Emissions(g/km)") \
    .show(20, truncate=False)
```

Vehicles with C02 Emissions greater than average:

Make	Model	Fuel Type	C02 Emissions(g/km)
ACURA	MDX 4WD	Z	255
ACURA	TL AWD	Z	255
ACURA	TL AWD	Z	267
ASTON MARTIN	DB9	Z	359
ASTON MARTIN	RAPIDE	Z	359
ASTON MARTIN	V8 VANTAGE	Z	338
ASTON MARTIN	V8 VANTAGE	Z	354
ASTON MARTIN	V8 VANTAGE S	Z	338
ASTON MARTIN	V8 VANTAGE S	Z	354
ASTON MARTIN	VANQUISH	Z	359
AUDI	A6 QUATTRO	Z	251
AUDI	A7 QUATTRO	Z	262
AUDI	A8	Z	258
AUDI	A8	Z	265
AUDI	A8L	Z	258
AUDI	A8L	Z	288
AUDI	A8L	Z	363

AUDI	Q5	Z	258
AUDI	Q7	Z	304
AUDI	Q7 TDI (modified)	D	290

only showing top 20 rows

#9. Increase CO2 by +10 for vehicles with Engine Size > 5.0L---

```
df_updated = df.withColumn("CO2 Emissions(g/km)",
                           col("CO2 Emissions(g/km)") +
                           (10 * (col("Engine Size(L)") >
5.0)).cast("int")))
```

```
print("Updated CO2 values for vehicles with Engine Size > 5.0L:")
df_updated.filter(col("Engine Size(L)") > 5.0).select("Make", "Engine
Size(L)", "CO2 Emissions(g/km)").show(20, truncate=False)
```

Updated CO2 values for vehicles with Engine Size > 5.0L:

Make	Engine Size(L)	CO2 Emissions(g/km)
ASTON MARTIN	5.9	369
ASTON MARTIN	5.9	369
ASTON MARTIN	5.9	369
AUDI	6.3	373
AUDI	5.2	357
AUDI	5.2	419
AUDI	5.2	357
AUDI	5.2	419
BENTLEY	6.0	366
BENTLEY	6.0	390
BENTLEY	6.0	390
BENTLEY	6.0	398
BENTLEY	6.8	447
BMW	6.0	366
CADILLAC	6.2	392
CADILLAC	6.2	350
CADILLAC	6.2	392
CADILLAC	6.2	350
CADILLAC	6.2	392
CADILLAC	6.2	350

only showing top 20 rows

#10. Save vehicles with CO2 > 300 into CSV-----

```
import os, shutil
```

Change paths to your location

```
output_dir = r"C:\Users\sathw\High_CO2_Vehicles_temp"
```

```

final_file = r"C:\Users\sathw\High_CO2_Vehicles.csv"
high_co2_df = df.filter(col("C02 Emissions(g/km)") > 300)
try:
    # Write Spark output to temp folder
    high_co2_df.coalesce(1) \
        .write.option("header", True) \
        .mode("overwrite") \
        .csv(output_dir)

    # Find the part file Spark created
    for file in os.listdir(output_dir):
        if file.startswith("part-") and file.endswith(".csv"):
            part_file = os.path.join(output_dir, file)
            # Move and rename to desired file
            shutil.move(part_file, final_file)
            break

    # Clean up temp directory
    shutil.rmtree(output_dir)

    print(f" Saved vehicles with C02 > 300 into: {final_file}")
except Exception as e:
    print("⚠ Spark CSV write failed, using Pandas fallback...", e)
    high_co2_df.toPandas().to_csv(final_file, index=False)
    print(" Saved vehicles with C02 > 300 using Pandas fallback.")

 Saved vehicles with C02 > 300 into: C:\Users\sathw\
High_CO2_Vehicles.csv

```

Conclusion

--Overall, the dataset portrays a diverse Canadian automotive market with a wide mix of vehicles across size, power, and efficiency. The absence of missing data makes this dataset reliable for environmental policy analysis, automotive research, and predictive modeling of emissions. It highlights key industry trends: strong SUV and pickup representation, dominance of gasoline fuel types, and the growing presence of fuel-efficient models.